

MP2-Passenger_Flight

November 16, 2023

```
[ ]: # Dy, Alwyn  
     # del Castillo, Kyle Adrian  
     # Lardizabal, Vladimir  
     # Quilario, Jose Angelo Virgil  
     # Segovia, Ethan
```

1 Instructions:

1. Create an overview of the problem being solved, e.g., what was the story behind the collection of the data, description of the attributes/features used, etc.
2. (Data Preprocessing and Exploratory Analysis) Present descriptive statistics as applicable (e.g., distribution, central tendency, variability) of the data before training the models. Clean the data if there are missing values, etc. You may perform feature engineering (i.e., creating new features out of the given features), but be sure to document your justifications.
3. Split your data into proportions of 70% training set and 30% testing set.
4. Train the following models: (a) logistic regression classifier and (b) naive Bayes classifier on the dataset.
5. Evaluate the performance of the trained model. You may use additional performance measures if you want, but for now I will only require the calculation of the accuracy. The accuracy measures the fraction of correct classifications. With this, you need to generate the confusion matrix. You may read this if you haven't encountered this concept before: <https://www.sciencedirect.com/topics/engineering/confusion-matrix#:~:text=A%20confusion%20matrix%20represents%20the,by%20model%20as%20other%20class>. Remember to compute this matrix from the test set (not the training set).

```
[31]: import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt  
import warnings  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import classification_report, confusion_matrix  
from sklearn.preprocessing import StandardScaler  
from sklearn.linear_model import LogisticRegression  
from sklearn.naive_bayes import GaussianNB  
from sklearn.metrics import accuracy_score
```

```
[32]: warnings.filterwarnings("ignore")
```

2 4. Passenger Flight

URL: <https://www.kaggle.com/datasets/teejmahal20/airline-passenger-satisfaction/data>

```
[33]: dataset = pd.read_csv("passenger_flight.csv")
np.random.seed(123)
dataset = dataset.sample(frac=1).reset_index(drop=True)
display(dataset.head(5))
display(dataset.tail(5))
display(dataset.info())
```

	Gender	Customer Type	Age	Type of Travel	Class	Flight Distance \
0	1	1	53	0	1	946
1	1	1	55	1	1	1620
2	1	1	29	1	1	1400
3	1	1	49	1	1	850
4	1	1	42	1	1	3535

	Inflight wifi service	Departure/Arrival time convenient \
0	1	1
1	3	3
2	5	2
3	5	5
4	5	5

	Ease of Online booking	Gate location ...	Inflight entertainment \
0	1	2 ...	2
1	3	3 ...	4
2	2	2 ...	5
3	5	5 ...	2
4	1	5 ...	5

	On-board service	Leg room service	Baggage handling	Checkin service \
0	2	5	2	4
1	4	4	4	3
2	4	2	2	5
3	2	3	2	5
4	5	5	5	3

	Inflight service	Cleanliness	Departure Delay in Minutes \
0	3	2	0
1	4	4	0
2	2	5	0
3	2	3	0
4	5	3	5

	Arrival Delay in Minutes	satisfaction
0	0.0	0
1	0.0	1
2	0.0	1
3	0.0	1
4	0.0	1

[5 rows x 23 columns]

	Gender	Customer Type	Age	Type of Travel	Class	Flight Distance	\
25971	0	1	44	1	1	3540	
25972	1	0	33	1	0	422	
25973	1	1	64	0	0	163	
25974	0	1	55	0	1	261	
25975	1	1	61	1	1	189	

	Inflight wifi service	Departure/Arrival time convenient	\
25971	1	1	
25972	1	1	
25973	3	5	
25974	3	2	
25975	4	4	

	Ease of Online booking	Gate location	...	Inflight entertainment	\
25971	1	1	...	5	
25972	1	3	...	5	
25973	3	2	...	2	
25974	0	3	...	3	
25975	4	4	...	3	

	On-board service	Leg room service	Baggage handling	Checkin service	\
25971	5	5	5	1	
25972	1	2	3	3	
25973	4	4	4	3	
25974	3	0	3	1	
25975	3	3	4	2	

	Inflight service	Cleanliness	Departure Delay in Minutes	\
25971	5	4	0	
25972	4	5	0	
25973	4	2	0	
25974	3	1	0	
25975	3	2	72	

	Arrival Delay in Minutes	satisfaction
25971	0.0	1
25972	0.0	0

```

25973          9.0          0
25974          0.0          0
25975         70.0          0

```

```
[5 rows x 23 columns]
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 25976 entries, 0 to 25975
```

```
Data columns (total 23 columns):
```

#	Column	Non-Null Count	Dtype
0	Gender	25976 non-null	int64
1	Customer Type	25976 non-null	int64
2	Age	25976 non-null	int64
3	Type of Travel	25976 non-null	int64
4	Class	25976 non-null	int64
5	Flight Distance	25976 non-null	int64
6	Inflight wifi service	25976 non-null	int64
7	Departure/Arrival time convenient	25976 non-null	int64
8	Ease of Online booking	25976 non-null	int64
9	Gate location	25976 non-null	int64
10	Food and drink	25976 non-null	int64
11	Online boarding	25976 non-null	int64
12	Seat comfort	25976 non-null	int64
13	Inflight entertainment	25976 non-null	int64
14	On-board service	25976 non-null	int64
15	Leg room service	25976 non-null	int64
16	Baggage handling	25976 non-null	int64
17	Checkin service	25976 non-null	int64
18	Inflight service	25976 non-null	int64
19	Cleanliness	25976 non-null	int64
20	Departure Delay in Minutes	25976 non-null	int64
21	Arrival Delay in Minutes	25893 non-null	float64
22	satisfaction	25976 non-null	int64

```
dtypes: float64(1), int64(22)
```

```
memory usage: 4.6 MB
```

```
None
```

2.1 Data Preprocessing

- Features will be uniformed
- Whitespaces in features will be replaced by '_'
- If there are any missing values, it will be replaced by the mean of its corresponding feature.

```

[34]: # Column names cleaning
dataset.columns = [c.replace(' ', '_') for c in dataset.columns]

# Find any missing values

```

```
def calculate_missing_values(data):
    total_missing = data.isnull().sum()

    missing_data = pd.DataFrame({
        'Total Missing': total_missing,
    })
    return missing_data

calculate_missing_values(dataset)
```

```
[34]:
```

	Total Missing
Gender	0
Customer_Type	0
Age	0
Type_of_Travel	0
Class	0
Flight_Distance	0
Inflight_wifi_service	0
Departure/Arrival_time_convenient	0
Ease_of_Online_booking	0
Gate_location	0
Food_and_drink	0
Online_boarding	0
Seat_comfort	0
Inflight_entertainment	0
On-board_service	0
Leg_room_service	0
Baggage_handling	0
Checkin_service	0
Inflight_service	0
Cleanliness	0
Departure_Delay_in_Minutes	0
Arrival_Delay_in_Minutes	83
satisfaction	0

```
[55]: # Replace missing values with its column's mean
dataset['Arrival_Delay_in_Minutes'] = dataset['Arrival_Delay_in_Minutes'].
    ↪ fillna(dataset['Arrival_Delay_in_Minutes'].mean())
calculate_missing_values(dataset)
```

```
[55]:
```

	Total Missing
Gender	0
Customer_Type	0
Age	0
Type_of_Travel	0
Class	0
Flight_Distance	0

Inflight_wifi_service	0
Departure/Arrival_time_convenient	0
Ease_of_Online_booking	0
Gate_location	0
Food_and_drink	0
Online_boarding	0
Seat_comfort	0
Inflight_entertainment	0
On-board_service	0
Leg_room_service	0
Baggage_handling	0
Checkin_service	0
Inflight_service	0
Cleanliness	0
Departure_Delay_in_Minutes	0
Arrival_Delay_in_Minutes	0
satisfaction	0

2.2 Exploratory Analysis

- Here we try to gain some basic understanding of our dataset

```
[39]: dataset.describe()
```

```
[39]:
```

	Gender	Customer_Type	Age	Type_of_Travel	\
count	25976.000000	25976.000000	25976.000000	25976.000000	
mean	0.492917	0.815253	39.620958	0.694410	
std	0.499959	0.388100	15.135685	0.460666	
min	0.000000	0.000000	7.000000	0.000000	
25%	0.000000	1.000000	27.000000	0.000000	
50%	0.000000	1.000000	40.000000	1.000000	
75%	1.000000	1.000000	51.000000	1.000000	
max	1.000000	1.000000	85.000000	1.000000	

	Class	Flight_Distance	Inflight_wifi_service	\
count	25976.000000	25976.000000	25976.000000	
mean	0.554820	1193.788459	2.724746	
std	0.496995	998.683999	1.335384	
min	0.000000	31.000000	0.000000	
25%	0.000000	414.000000	2.000000	
50%	1.000000	849.000000	3.000000	
75%	1.000000	1744.000000	4.000000	
max	1.000000	4983.000000	5.000000	

	Departure/Arrival_time_convenient	Ease_of_Online_booking	\
count	25976.000000	25976.000000	
mean	3.046812	2.756775	

std	1.533371	1.412951
min	0.000000	0.000000
25%	2.000000	2.000000
50%	3.000000	3.000000
75%	4.000000	4.000000
max	5.000000	5.000000

	Gate_location ...	Inflight_entertainment	On-board_service \
count	25976.000000 ...	25976.000000	25976.000000
mean	2.977094 ...	3.357753	3.385664
std	1.282133 ...	1.338299	1.282088
min	1.000000 ...	0.000000	0.000000
25%	2.000000 ...	2.000000	2.000000
50%	3.000000 ...	4.000000	4.000000
75%	4.000000 ...	4.000000	4.000000
max	5.000000 ...	5.000000	5.000000

	Leg_room_service	Baggage_handling	Checkin_service	Inflight_service \
count	25976.000000	25976.000000	25976.000000	25976.000000
mean	3.350169	3.633238	3.314175	3.649253
std	1.318862	1.176525	1.269332	1.180681
min	0.000000	1.000000	1.000000	0.000000
25%	2.000000	3.000000	3.000000	3.000000
50%	4.000000	4.000000	3.000000	4.000000
75%	4.000000	5.000000	4.000000	5.000000
max	5.000000	5.000000	5.000000	5.000000

	Cleanliness	Departure_Delay_in_Minutes	Arrival_Delay_in_Minutes \
count	25976.000000	25976.000000	25893.000000
mean	3.286226	14.30609	14.740857
std	1.319330	37.42316	37.517539
min	0.000000	0.000000	0.000000
25%	2.000000	0.000000	0.000000
50%	3.000000	0.000000	0.000000
75%	4.000000	12.000000	13.000000
max	5.000000	1128.000000	1115.000000


	satisfaction
count	25976.000000
mean	0.438982
std	0.496272
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	1.000000

[8 rows x 23 columns]

2.2.1 We will do some more visualization of the data to understand each feature


```
[40]: dataset['satisfaction'] = dataset['satisfaction'].astype(str)
```

```
[48]: # Gender vs Satisfaction
with sns.axes_style(style='ticks'):
    g = sns.catplot(x="satisfaction", col="Gender", col_wrap=2, data=dataset,
    ↪kind="count", height=3, aspect=1.0, palette="bright", hue="satisfaction",
    ↪legend=False)
    g.add_legend()
```




It can be observed that the distribution of satisfied and dissatisfied are quite the same between both genders. Dissatisfied customers are higher in number compared to the satisfied customers.

```
[49]: # Customer Type vs Satisfaction
with sns.axes_style(style='ticks'):
    g = sns.catplot(x="satisfaction", col="Customer_Type", col_wrap=2,
    ↪data=dataset, kind="count", height=3, aspect=1.0, palette="bright",
    ↪hue="satisfaction", legend=False)
```


There is more loyal customers compared to the disloyal customers, which gives us the insight that most of the people in the sample are people that have already flown before. We can also see that the proportionality of satisfaction to customer type is significantly high.

```
[58]: # Age vs Satisfaction
with sns.axes_style('white'):
    g = sns.catplot(x="Age", data=dataset, aspect=3.0, kind='count',
    ↪ hue='satisfaction', order=range(5, 80), palette="bright")
    g.set_axis_labels('Age', 'Count')
    g.set(title='Age vs Passenger Satisfaction')
```



```
[59]: # Type of Travel, Class vs Satisfaction
with sns.axes_style('white'):
    g = sns.catplot(x="Type_of_Travel", hue="satisfaction", col="Class",
    ↪ data=dataset, kind="count", height=5, aspect=.8, palette="bright")
```

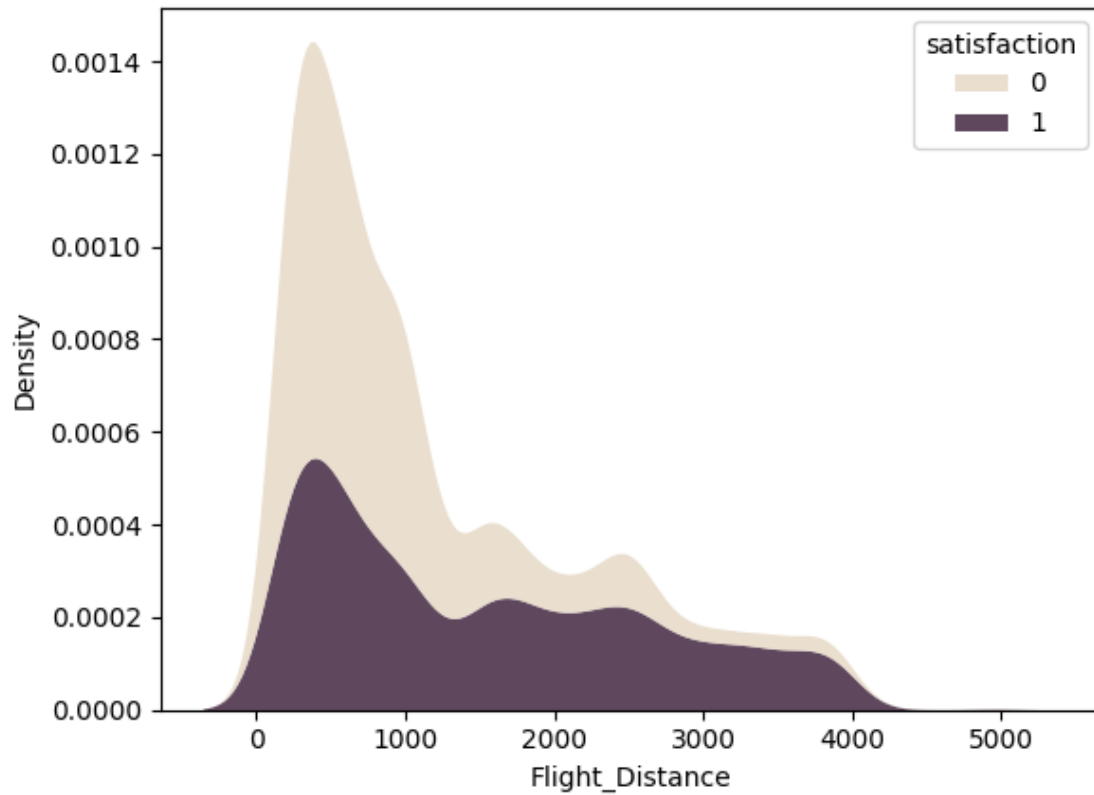
```
g.set_titles("Class {col_name} | Type of Travel: Personal")
g.axes[0, 1].set_title("Class 1 | Type of Travel: Business")
```



Personal Travel has more dissatisfied customers compared to business travels. Business travels has more satisfied customers.


```
[60]: # Flight Distance vs Satisfaction
sns.kdeplot(data = dataset, x = "Flight_Distance", hue = "satisfaction" , shade_
↪ = True, palette="ch:.25", multiple="stack", fill=True, common_norm=False, alpha=.
↪ 8, linewidth=0)
```

```
[60]: <Axes: xlabel='Flight_Distance', ylabel='Density'>
```



```
[61]: # Delay
sns.scatterplot(x = 'Departure_Delay_in_Minutes', y = 'Arrival_Delay_in_Minutes', data = dataset, palette="ch:.25")
```


[61]: <Axes: xlabel='Departure_Delay_in_Minutes', ylabel='Arrival_Delay_in_Minutes'>



```
[62]: # Satisfaction
satisfaction_counts = dataset['satisfaction'].value_counts()

# Plotting the pie chart
plt.figure(figsize=(6, 6))
plt.pie(satisfaction_counts, labels=satisfaction_counts.index, autopct='%1.1f%%', colors=['lightcoral', 'lightblue'])
plt.title('Distribution of Satisfaction')
```

```
[62]: Text(0.5, 1.0, 'Distribution of Satisfaction')
```



Gender: It can be observed that the distribution of satisfied and dissatisfied are quite the same between both genders. **Dissatisfied customers** are higher in number compared to the satisfied customers. **Customer Type:** Loyal customers are higher in number compared to the disloyal customers. **Age:** In age groups 39-60, there are more satisfied customers than the dissatisfied customers. **Type of Travel & Class:** Personal Travel has more dissatisfied customers compared to business travels. Business travels has more satisfied customers. **Satisfaction:** There are more dissatisfied customers than satisfied customers. It is worth note taking that more than half are dissatisfied customers.

2.3 Detection of Outliers and Removal

Interquartile Range (IQR) will be used to detect outliers. Outliers will be removed from the dataset.

```
[11]: Q1 = dataset.quantile(0.25)
      Q3 = dataset.quantile(0.75)
      IQR = Q3 - Q1
      print(IQR)
```

Gender	1.0
Customer Type	0.0
Age	24.0
Type of Travel	1.0
Class	1.0
Flight Distance	1330.0
Inflight wifi service	2.0
Departure/Arrival time convenient	2.0
Ease of Online booking	2.0
Gate location	2.0
Food and drink	2.0
Online boarding	2.0
Seat comfort	3.0
Inflight entertainment	2.0
On-board service	2.0
Leg room service	2.0
Baggage handling	2.0
Checkin service	1.0
Inflight service	2.0
Cleanliness	2.0
Departure Delay in Minutes	12.0
Arrival Delay in Minutes	13.0
satisfaction	1.0

dtype: float64


```
[24]: # Remove Outliers
cleaned_ds = dataset[~((dataset < (Q1 - 1.5 * IQR)) |(dataset > (Q3 + 1.5 *
↪IQR))).any(axis=1)]
print(f'Number of rows before removal: {len(dataset)}')
print(f'Number of rows before removal: {len(cleaned_ds)}')
```

Number of rows before removal: 25976
Number of rows before removal: 25976

2.4 Features Correlation

```
[78]: correlation=cleaned_ds.corr()
plt.figure(figsize=(14,8))
sns.heatmap(correlation,annot=True,fmt='.2f',annot_kws={'size':↪
↪10},linewidths=0.5,cmap='Blues')
plt.title("Data correlations")
```

```
[78]: Text(0.5, 1.0, 'Data correlations')
```



2.5 Splitting of Data

```
[25]: x = cleaned_ds.drop('satisfaction', axis=1)
      y = cleaned_ds['satisfaction']

      x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.3,
      random_state=42)
      print("Training set shape:", x_train.shape, y_train.shape)
      print("Testing set shape:", x_test.shape, y_test.shape)
```

Training set shape: (18183, 22) (18183,)

Testing set shape: (7793, 22) (7793,)

2.6 Confusion Matrix

```
[26]: def make_confusion_matrix(cf,
      group_names=None,
      categories='auto',
      count=True,
      percent=True,
      cbar=True,
      xyticks=True,
```

```

        xyplotlabels=True,
        sum_stats=True,
        figsize=None,
        cmap='Blues',
        title=None):

    blanks = ['' for i in range(cf.size)]

    if group_names and len(group_names)==cf.size:
        group_labels = ["{}\n".format(value) for value in group_names]
    else:
        group_labels = blanks

    if count:
        group_counts = ["{0:0.0f}\n".format(value) for value in cf.flatten()]
    else:
        group_counts = blanks

    if percent:
        group_percentages = ["{0:.2%}".format(value) for value in cf.flatten()/
↪np.sum(cf)]
    else:
        group_percentages = blanks

    box_labels = [f"{v1}-{v2}-{v3}".strip() for v1, v2, v3 in_
↪zip(group_labels,group_counts,group_percentages)]
    box_labels = np.asarray(box_labels).reshape(cf.shape[0],cf.shape[1])

    if sum_stats:
        accuracy = np.trace(cf) / float(np.sum(cf))

        if len(cf)==2:
            #Metrics for Binary Confusion Matrices
            precision = cf[1,1] / sum(cf[:,1])
            recall = cf[1,1] / sum(cf[1,:])
            f1_score = 2*precision*recall / (precision + recall)
            stats_text = "\n\nAccuracy={:0.3f}\nPrecision={:0.3f}\nRecall={:0.
↪3f}\nF1 Score={:0.3f}".format(
                accuracy,precision,recall,f1_score)
        else:
            stats_text = "\n\nAccuracy={:0.3f}".format(accuracy)
    else:
        stats_text = ""

    if figsize==None:
        figsize = plt.rcParams.get('figure.figsize')

```



```

    if xyticks==False:
        categories=False

    plt.figure(figsize=figsize)
    sns.
    ↪heatmap(cf,annot=box_labels,fmt="",cmap=cmap,cbar=cbar,xticklabels=categories,yticklabels=c

    if xyplotlabels:
        plt.ylabel('True label')
        plt.xlabel('Predicted label' + stats_text)
    else:
        plt.xlabel(stats_text)

    if title:
        plt.title(title)

```

2.7 Logistic Regression

```

[15]: # Feature Scaling
sc = StandardScaler()
x_train = sc.fit_transform(x_train)
x_test = sc.transform(x_test)

```

```

[138]: class LogisticRegressionScratch:
    def __init__(self, learning_rate=0.01, num_iterations=300000,
    ↪fit_intercept=True, verbose=False):
        self.learning_rate = learning_rate
        self.num_iterations = num_iterations
        self.fit_intercept = fit_intercept
        self.verbose = verbose

    def __b_intercept(self, X):
        intercept = np.ones((X.shape[0], 1))
        return np.concatenate((intercept, X), axis=1)

    def __sigmoid_function(self, z):
        return 1 / (1 + np.exp(-z))

    def __loss(self, yp, y):
        return (-y * np.log(yp) - (1 - y) * np.log(1 - yp)).mean()

    def fit(self, X, y):
        if self.fit_intercept:
            X = self.__b_intercept(X)
        self.W = np.zeros(X.shape[1])

```

```

    for i in range(self.num_iterations):

        z = np.dot(X, self.W)
        yp = self.__sigmoid_function(z)
        gradient = np.dot(X.T, (yp - y)) / y.size

        self.W -= self.learning_rate * gradient

        z = np.dot(X, self.W)
        yp = self.__sigmoid_function(z)

        loss = self.__loss(yp, y)

        if(self.verbose ==True and i % 10000 == 0):
            print(f'loss: {loss} \t')

    def predict_prob(self, X):
        if self.fit_intercept:
            X = self.__b_intercept(X)
        return self.__sigmoid_function(np.dot(X, self.W))

    def predict(self, X):
        return self.predict_prob(X).round()

```

```

[139]: log_res = LogisticRegressionScratch(learning_rate=0.1, num_iterations=300000)
log_res.fit(x_train, y_train)

```


```

[156]: print(classification_report(y_test, log_res.predict(x_test)))

cm = confusion_matrix(y_test, log_res.predict(x_test))
make_confusion_matrix(cm)

```

	precision	recall	f1-score	support
0	0.87	0.91	0.89	4416
1	0.87	0.82	0.85	3377
accuracy			0.87	7793
macro avg	0.87	0.86	0.87	7793
weighted avg	0.87	0.87	0.87	7793



2.7.1 Using the Library

```
[154]: lr = LogisticRegression()
lr.fit(x_train, y_train)
```


```
[154]: LogisticRegression()
```

```
[155]: y_pred = lr.predict(x_test)
print(classification_report(y_test, y_pred))

cm = confusion_matrix(y_test, y_pred)
make_confusion_matrix(cm)
```

	precision	recall	f1-score	support
0	0.87	0.91	0.89	4416
1	0.87	0.82	0.85	3377

accuracy			0.87	7793
macro avg	0.87	0.86	0.87	7793
weighted avg	0.87	0.87	0.87	7793



Accuracy=0.870
 Precision=0.871
 Recall=0.822
 F1 Score=0.846

2.8 Naive Bayes

2.8.1 Using Builtin Lib


```
[146]: model = GaussianNB()
      model.fit(x_train, y_train)
```

```
[146]: GaussianNB()
```

```
[157]: y_pred = model.predict(x_test)
```

```
print(classification_report(y_test, y_pred))
cm = confusion_matrix(y_test, y_pred)
make_confusion_matrix(cm)
```

	precision	recall	f1-score	support
0	0.86	0.90	0.88	4416
1	0.87	0.81	0.84	3377
accuracy			0.86	7793
macro avg	0.86	0.86	0.86	7793
weighted avg	0.86	0.86	0.86	7793



Accuracy=0.865
Precision=0.865
Recall=0.815
F1 Score=0.839